

Neural Network Assignment

Group of Aline, Norma, Samuel and Elise

Introduction

Aim: create a NN model or use a pretrained one and explore its behavior.

- Dataset: MNIST image data
- Approaches:
 - Paper PDF (Aline)
 - Google Gemini (Elise)
 - Chat GPT (Sam)
 - Claude (Norma)
- Prompt for (2,3,4):

I need to create a neural network model in Python for my ^
studies, using the MNIST image data set, to recognise
handwritten single digits. Ask me any follow-up
questions so you have enough information to create a
working code

- Organisation: 4 online meetings + WhatsApp group

Only Book (Aline)

```
import pandas as pd  
  
obj = pd.read_pickle(r'/Users/socialmediaappleid/Desktop/CAS AICP/Python/neural-networks-and-deep-learning-master/datasets/mnist/mnist.pkl')  
Python
```

```
NameError Traceback (most recent call last)  
Cell In[5], line 1  
----> 1 net = network.Network([784, 100, 10])  
      2 net.SGD(training_data, 30, 10, 0.001, test_data=test_data)  
  
NameError: name 'network' is not defined
```

ModuleNotFoundError: No module named 'cPickle'

```
NameError Traceback (most recent call last)  
Cell In[6], line 1  
----> 1 net.SGD(training_data, 30, 10, 3.0, test_data=test_data)  
  
NameError: name 'training_data' is not defined
```

Generate + Code + Markdown

```
NameError Traceback (most recent call last)  
Cell In[1], line 2  
      1 selection_list = []  
----> 2 get_selection_indices(x=embedding[:, 0], y=embedding[:, 1], selection=selection)  
  
NameError: name 'get_selection_indices' is not defined
```

```
ModuleNotFoundError Traceback (most recent call last)  
Cell In[2], line 1  
----> 1 import mnist_loader  
      2 training_data, validation_data, test_data = \  
      3 mnist_loader.load_data_wrapper()  
  
ModuleNotFoundError: No module named 'mnist_loader'
```

```
File /usr/lib/python3.12/random.py:357, in Random.shuffle(self, x)  
 354 for i in reversed(range(1, len(x))):  
 355     # pick an element in x[:i+1] with which to exchange x[i]  
 356     j = randbelow(i + 1)  
--> 357     x[i], x[j] = x[j], x[i]  
  
TypeError: 'DataLoader' object is not subscriptable
```

```
print "Training with expanded data, %s neurons in the FC layer, run num %s" % (n, j)
```

```

class Network(object):
    def __init__(self, sizes):
        """
        The list `sizes` contains the number of neurons in the
        respective layers of the network. For example, if the list
        was [2, 3, 1] then it would be a three-layer network, with the
        first layer containing 2 neurons, the second layer 3 neurons,
        and the third layer 1 neuron.
        """
        self.num_layers = len(sizes)
        self.sizes = sizes
        # biases for all layers except input
        self.biases = [np.random.randn(y, 1) for y in sizes[1:]]
        # weights between layers
        self.weights = [np.random.randn(y, x) for x, y in zip(sizes[:-1], sizes[1:])]

    def feedforward(self, a):
        """
        Return the output of the network if `a` is input.
        """
        for b, w in zip(self.biases, self.weights):
            a = sigmoid(np.dot(w, a) + b)
        return a

    def SGD(self, training_data, epochs, mini_batch_size, eta, test_data=None):
        """
        Train the neural network using mini-batch stochastic gradient descent.
        `training_data` is a list of (x, y) tuples.
        If `test_data` is provided, the network is evaluated after each epoch.
        """
        if test_data:
            n_test = len(test_data)
        # ensure we can shuffle (in case an iterator is passed in)
        training_data = list(training_data)
        n = len(training_data)
        for j in range(epochs):
            random.shuffle(training_data)
            mini_batches = [
                training_data[k:k + mini_batch_size]
                for k in range(0, n, mini_batch_size)]

```

Book

```

class MyModel(nn.Module):
    def __init__(self, n_input, n_hiddens, n_output):
        super(MyModel, self).__init__()

        self.ls = []
        n_prev = n_input
        for i, n_out in enumerate(n_hiddens):
            l = nn.Linear(n_prev, n_out) # for hidden layer we create a linear projection from n_prev features to n_out
            n_prev = n_out
            self.add_module(f'lin_{i}_{n_out}', l)
            self.ls.append(l)

        self.lout = nn.Linear(n_prev, n_output) # also we need the output layer

    def forward(self, x):
        h = x
        for li in self.ls: # for each layer we apply the linear projection and the activation function (ReLU)
            h = li(h)
            h = torch.relu(h)

        logits = self.lout(h)
        # Apply softmax activation per row, to get the class pseudoprobabilities
        probs = F.softmax(logits, dim=1)

        # Prediction: argmax for classification
        pred = torch.argmax(probs, dim=1) # find the element with highest value in each row

        return logits, probs, pred

```

✓ 0.0s

Python

Tutorial

```
# Define a transform to normalize the data
transform = transforms.Compose([transforms.ToTensor(),
                               transforms.Normalize((0.5,), (0.5,)),
                               transforms.Lambda(lambda x: torch.flatten(x))])

# Download and load the training data
trainset = datasets.FashionMNIST('~./pytorch/FMNIST_data/',
                                  download=True, train=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=64, shuffle=True)

# Download and load the test data
testset = datasets.FashionMNIST('~./pytorch/FMNIST_data/',
                                 download=True, train=False, transform=transform)
testloader = DataLoader(testset, batch_size=64, shuffle=True)

# Print shapes of the datasets
print('Train dataset shape:', len(trainset), 'total images and labels')
print('Test dataset shape:', len(testset), 'total images and labels')
```

My Learnings

- I need new Materials and/or the help of friends/Chatgpt
- >>> means it has to go into the Terminal
- In Python 3, cPickle is called Pickle

Starting with Gemini

- Definition with Gemini: deep learning framework (Pytorch), NN architecture (MLP), my experience level with Python (Beginner), specific requirements for the model and how the code should be structured (Jupyter Notebook)
- Approach: using Gemini as my really smart friend and challenge him with the differences that I could see with the exercice during the course
- Play and advise with different parameters: n_hiddens, num_epochs, lr and batch_size

The Neural Network Model

- Input Layer: 784 numbers from 28x28pixels
- Hidden layers: final model is with 128 neurons
- Output layer: 10 (number from 0 to 9)

Training the model

- Data Loading with DataLoader
- Loss function: Cross-Entropy Loss
- Optimiser: Adam optimiser
- Epochs: 20

Testing my own handwriting


- Integration of my own digits and converted into the same format as the MNIST data
- Image added to the same directory as my notebook
- **Preprocess the image**
- Get a prediction
- **Change parameters:** hidden layers (128), lr (0,005), epoch (20)
- Try with a new digit

```
# Load the image and convert it to grayscale
img = Image.open(image_path).convert('L')

# Resize to 28x28 pixels
img = img.resize((28, 28), Image.LANCZOS)
.

# Convert to a NumPy array
img_array = np.array(img)

# Invert the image if needed (white background, black digit)
if np.mean(img_array) > 128:
    img_array = 255 - img_array
```



Learnings

1. A model might memorize the training data well but struggle to generalize to new, unseen data
2. The diversity of its training data is so important : MNIST specific handwriting style
3. Hyperparameters will optimise the model's performance but still can fail with new unseen data

Building a handwritten digit classifier with Chat GPT (Sam)

- I was tasked with creating a model using **ChatGPT** as my coding assistant.
- Chat GPT setup a workflow in PyTorch, running on my MacBook M2 GPU.

Approach:

- Baseline CNN:** simple convolutional model for a strong starting point.
- Small ResNet:** deeper model with skip connections, aimed at higher accuracy.

Goal 1: Train, evaluate, and compare both models on the MNIST dataset.

Goal 2: Create additional code to ingest my own handwritten digits and recognize them.

```
# ----- 7) Models -----
class SimpleCNN(nn.Module):
    def __init__(self, num_classes=10):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(1, 32, 3, padding=1), nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, 3, padding=1), nn.ReLU(inplace=True),
            nn.MaxPool2d(2), # 14x14
            nn.Conv2d(32, 64, 3, padding=1), nn.ReLU(inplace=True),
            nn.Conv2d(64, 64, 3, padding=1), nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64 * 7 * 7, 128), nn.ReLU(inplace=True),
            nn.Linear(128, num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        return self.classifier(x)

# defines a simple convolutional neural network (CNN) class for image classification
# initializes the SimpleCNN class with a specified number of output classes (default is 10 for MNIST digits 0-9)
# initializes the parent class (nn.Module)
# defines a sequential container for the feature extraction layers of the CNN, meaning that the layers will be applied in order
# first convolutional layer with 1 input channel (grayscale), 32 output channels, kernel size of 3, and ReLU activation
# In a convolutional layer, the kernel (or filter) size determines the dimensions of the sliding window that
# moves over the input image to extract features.
# A kernel size of 3 means a 3x3 pixels filter is used to scan over the input image.
# max pooling layer with a kernel size of 2, which reduces the spatial dimensions by half (from 28x28 to 14x14).
# A Max Pooling layer is a downsampling operation that
# reduces the spatial dimensions (height and width) of the input feature maps.
# second convolutional layer with 32 input channels, 64 output channels, kernel size of 3, and ReLU activation
# third convolutional layer with 64 input channels, 64 output channels, kernel size of 3, and ReLU activation
# 7x7 # reduces spatial dimensions to 7x7
# defines a sequential container for the classification layers of the CNN
# flattens the 2D feature maps into a 1D vector
# fully connected layer with 64x7x7 input features and 128 output features, followed by ReLU activation
# final fully connected layer with 128 input features and num_classes output features (logits for each class)
```

Screenshot of the Convolutional Neural Network model (CNN)

```
class BasicBlock(nn.Module):
    # defines a basic residual block class for a ResNet architecture.
    # Residual Networks, a type of neural network that uses skip connections
    # to facilitate gradient flow during training, helping to mitigate the vanishing gradient problem in deep networks.

    def __init__(self, in_ch, out_ch, stride=1):
        super().__init__()
        self.conv1 = nn.Conv2d(in_ch, out_ch, 3, stride=stride, padding=1, bias=False) # first convolutional layer with specified input/output channels, kernel size of 3, stride,
        self.bn1 = nn.BatchNorm2d(out_ch) # batch normalization layer for the output of the first convolutional layer, meaning that it normalizes
                                        # the activation to improve training stability
        self.conv2 = nn.Conv2d(out_ch, out_ch, 3, padding=1, bias=False) # second convolutional layer with input channels as both input and output, kernel size of 3, padding of 1,
                                                                # and no bias
        self.bn2 = nn.BatchNorm2d(out_ch) # batch normalization layer for the output of the second convolutional layer


        self.downsample = None # initializes a downsampling layer as None, which will be used if the input and output dimensions do not match
        if stride > 1 or in_ch != out_ch: # checks if downsampling is needed based on stride or channel mismatch
            self.downsample = nn.Sequential( # defines a sequential container for the downsampling layer to match dimensions
                nn.Conv2d(in_ch, out_ch, 1, stride=stride, bias=False), # 1x1 convolutional layer to adjust the number of channels and spatial dimensions
                nn.BatchNorm2d(out_ch) # batch normalization layer to match the number of channels and spatial dimensions
            )

    def forward(self, x):
        identity = x # defines the forward pass of the BasicBlock class
        out = self.bn1(self.conv1(x)) # applies the first convolutional layer, followed by batch normalization and ReLU activation
        out = self.bn2(self.conv2(out)) # applies the second convolutional layer, followed by batch normalization
        if self.downsample is not None: # adds the identity connection from the input to the output
            identity = self.downsample(identity) # applies the downsampling layer to the identity
        out = F.relu(out + identity) # adds the identity (skip) connection to the output and applies ReLU activation
        return out # returns the output of the block
```



```
class SmallResNet(nn.Module):
    # defines a small ResNet architecture for image classification
    # A 3 layer ResNet with 32, 64, 128, one block per stage
    def __init__(self, num_classes=10):
        super().__init__()
        self.stem = nn.Sequential( # initializes the stem (or base) of the ResNet
            nn.Conv2d(1, 32, 3, padding=1, bias=False), # initial convolutional layer with 1 input channel (grayscale), 32 output channels, kernel size of 3, padding of 1, and no bias
            nn.BatchNorm2d(32), # batch normalization layer applied in-place to save memory
            nn.ReLU(inplace=True),
        )
        self.layer1 = BasicBlock(32, 32, stride=1) # 32x32 # first residual block with 32 input and output channels, stride of 1 (no downsampling)
        self.layer2 = BasicBlock(32, 64, stride=2) # 16x16 # second residual block with 32 input channels, 64 output channels, and stride of 2 (downsampling to reduce spatial dimensions)
        self.layer3 = BasicBlock(64, 128, stride=2) # 8x8 # third residual block with 64 input channels, 128 output channels, and stride of 2 (further downsampling)
        self.adaptiveAvgPooling1 = nn.AdaptiveAvgPooling2d(1) # adaptive average pooling layer that reduces the spatial dimensions of the feature maps of the hidden
        self.flatten = nn.Flatten() # flattens the 2D feature maps into a 1D vector
        self.linear1 = nn.Linear(128, num_classes) # final fully connected layer with 128 input features and num_classes output features (logits for each class)
```

Screenshots of the ResNet model



Overcoming Training Challenges


- The CNN worked well in the very first try. No code adjustments necessary.
- Initial training of ResNet failed (~10% accuracy). Chat GPT discovered that **EMA + BatchNorm** caused problems on Apple's MPS.
 - Solution: It suggested to turn off AMP and EMA for the ResNet. Then training "stabilized".

Results:

- Baseline CNN: 99.3% accuracy.
- Small ResNet: 99.6% accuracy (slightly better, as expected).

Key lesson from the ResNet fail:

Small details (like mixed precision, normalization layers) can make or break training stability.



Confusion Matrix Heatmap (Tries 1 to 3 from left to right)

== FINAL SUMMARY ==		
Model	Test Loss	Test Accuracy
Baseline CNN	0.0265	0.9927
Small ResNet + LS + EMA	28.271	0.0974

Summary 1st try (untouched ResNet)

== FINAL SUMMARY ==		
Model	Test Loss	Test Accuracy
Baseline CNN	0.0265	0.9927
Small ResNet + LS + EMA	3.3315	0.3345

Summary 2nd try (turned off AMP)

== FINAL SUMMARY ==		
Model	Test Loss	Test Accuracy
Baseline CNN	0.0265	0.9927
Small ResNet + LS + EMA	0.3036	0.9959

Summary 3rd try (turned off AMP & EMA)

Unanticipated code running errors:
Error 1: Using "Tabulate" without installing it
Error 2: Dataloader: Issue w/ VSCode/Notebooks

Testing with My Own Handwriting

- Chat GPT generated additional code to collected 10 of my own handwritten digits (0–9).
- It built a preprocessing pipeline to resize, invert, and normalize images to match MNIST style.

First attempt:



- all predictions for each of the 10 digits were “8” (showed preprocessing mismatch).

Second attempt:


- After increasing contrast in Photoshop: models reached ~70% accuracy on my digits.

Next suggested steps:


Chat GPT suggests to refine preprocessing code (force inversion, stronger contrast, stroke dilation) for better results.



```
== Results for Small ResNet ==
['file', 'true (from name)', 'pred', 'confidence']
['0.png', '0', 8, '0.98']
['1.png', '1', 8, '0.98']
['2.png', '2', 8, '0.98']
['3.png', '3', 8, '0.98']
['4.png', '4', 8, '0.98']
['5.png', '5', 8, '0.98']
['6.png', '6', 8, '0.98']
['7.png', '7', 8, '0.98']
['8.png', '8', 8, '0.98']
['9.png', '9', 8, '0.98']
Mini-set accuracy (from filenames): 1/10 = 0.100
```



```
== Results for Small ResNet ==
['file', 'true (from name)', 'pred', 'confidence']
['0.png', '0', 0, '0.85']
['1.png', '1', 1, '0.69']
['2.png', '2', 8, '0.43']
['3.png', '3', 3, '0.49']
['4.png', '4', 4, '0.79']
['5.png', '5', 1, '0.60']
['6.png', '6', 6, '0.90']
['7.png', '7', 7, '0.33']
['8.png', '8', 8, '0.97']
['9.png', '9', 1, '0.29']
Mini-set accuracy (from filenames): 7/10 = 0.700
```



```
== Results for Small ResNet ==
['file', 'true (from name)', 'pred', 'confidence']
['0.png', '0', 0, '0.87']
['1.png', '1', 4, '0.85']
['2.png', '2', 8, '0.48']
['3.png', '3', 8, '0.48']
['4.png', '4', 8, '0.39']
['5.png', '5', 5, '0.79']
['6.png', '6', 6, '0.33']
['7.png', '7', 8, '0.42']
['8.png', '8', 8, '0.98']
['9.png', '9', 8, '0.54']
Mini-set accuracy (from filenames): 4/10 = 0.400
```

MNIST Optimization and Modern Deep Learning Techniques with Claude

Feedforward Neural Network

Challenge: Outdated Book Example (2015)

Systematic MNIST Improvement

Original (2015)	Modern Version (2025)	Improvement
Sigmoid + SGD + MSE	ReLU + Adam + CrossEntropy	+247 predictions
95.01% (9501/10000)	97.48% (9748/10000)	(+2.47%)

python

```
# Original (outdated standards)
model = nn.Sequential([
    nn.Linear(784, 30),
    nn.Sigmoid(),                                # Vanishing Gradient Problem
    nn.Linear(30, 10),
    nn.Sigmoid()
])
optimizer = optim.SGD(lr=3.0)      # Fixed Learning Rate
criterion = nn.MSELoss()          # Suboptimal for classification
```

python

```
# Modern (current best practices)
model = nn.Sequential([
    nn.Linear(784, 128),
    nn.ReLU(),                                     # Solves Vanishing Gradient
    nn.Linear(128, 64),
    nn.ReLU(),
    nn.Linear(64, 10)
])
optimizer = optim.Adam(lr=0.001)    # Adaptive Learning Rate
criterion = nn.CrossEntropyLoss()  # Optimal for classification
```

Personalization to Own Handwriting

Model	MNIST Performance	Personal Handwriting (10 PNG files)
Original	95.01%	10% (1/10)
Improved	97.48%	30% (3/10)

Model	Personal Handwriting	Robustness
Fine-tuned	100% (10/10)	100% (Overfitting-Risk)
Augmented	100% (10/10)	90% on variations

- **Automatic extraction failed:** OpenCV recognized 0 valid digits
- **Large dataset impractical:** Writing hundreds of digits manually too time-consuming
- **Solution:** 10 handwritten PNG files (digit_0.png to digit_9.png)

python

```
# 10 Originalbilder → 210 Variationen
transforms = [
    transforms.RandomRotation(degrees=12),
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
    transforms.RandomAffine(degrees=0, scale=(0.8, 1.2))
]
```

Synthetic Multi-Digit Pictures

- **First evaluation (78.5% accuracy):** tested on *all* 200 images (training + test combined)
- **Second evaluation (2.5% accuracy):** tested only on 40 test images (20% split)

Why this happens:

- Dataset too small: 200 images are not enough for sequence-to-sequence learning
- No regularization: the model memorizes instead of learning
- Architecture issue: the current CNN→Dense setup is not suitable for variable sequence lengths
- **Actual test accuracy:** ~2.5% (on unseen data)
Apparent accuracy: ~78.5% (on training data)
Overfitting gap: ~76%

Evaluieren Multi-Digit Modell...			
True:	0	Pred:	0 ✓
True:	1	Pred:	9 ✗
True:	2	Pred:	0 ✗
True:	3	Pred:	0 ✗
True:	4	Pred:	0 ✗
True:	5	Pred:	0 ✗
True:	6	Pred:	0 ✗
True:	7	Pred:	0 ✗
True:	8	Pred:	7 ✗
True:	9	Pred:	0 ✗
MULTI-DIGIT ACCURACY RESULTS:			
Overall:	157/200	=	0.785
1-stellig:	1/10	=	0.100
2-stellig:	79/90	=	0.878
3-stellig:	77/100	=	0.770

Learnings

1. Best technique does not always mean improvement
2. Augmentation helps with generalizing and thus have a better accuracy
3. Even if its able to create new synthetic data, it still is not enough for creating meaningful results

Conclusion